

Machine Learning

Linear Regression

Professor: Cheolsoo Park



Acknowledgement

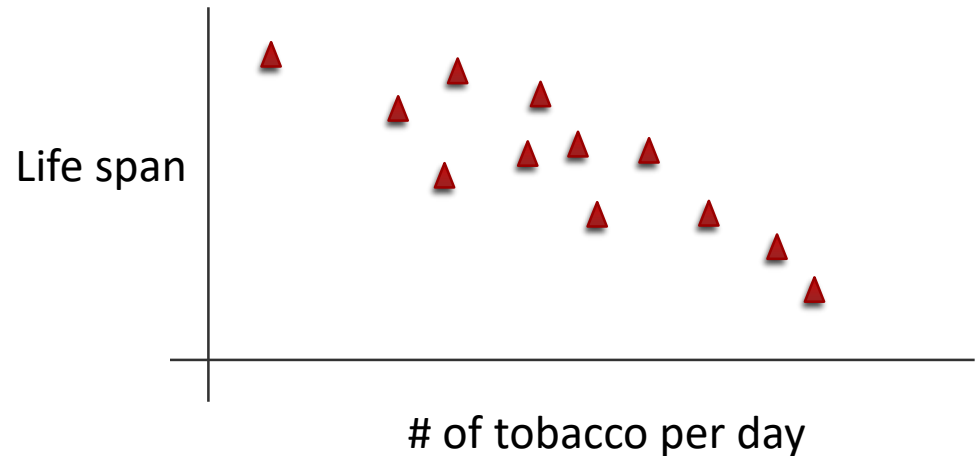


➤ Andrew Ng's ML class

➤ <https://www.coursera.org/learn/machine-learning>

One Variable Linear Regression

➤ Smoking vs lifespan




➤ Supervised Learning

- Classification : discrete output
- Regression : continuous value output



One Variable Linear Regression



➤ Linear regression with one variable

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

Cost function

➤ Parameter estimation in linear regression problem

➤ Hypothesis : $h_{\theta}(x) = \theta_0 + \theta_1 x$

, where θ_i ($i = 0$ and 1) are parameters

➤ Let's find the parameters, θ_i

➤ Cost function

$$J(\theta_0, \theta_1) = \frac{1}{2n} \sum_{k=1}^n (h_{\theta}(x^{(k)}) - y^{(k)})^2$$

➤ To estimate θ_0 and θ_1 , minimize $J(\theta_0, \theta_1)$

$$\underset{\theta_0, \theta_1}{\text{minimize}} J(\theta_0, \theta_1)$$

Cost Function

- Linear regression hypothesis

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

- Cost function

$$J(\theta_0, \theta_1) = \frac{1}{2n} \sum_{k=1}^n (h_{\theta}(x^{(k)}) - y^{(k)})^2$$

- To estimate θ_0 and θ_1 , minimize $J(\theta_0, \theta_1)$

$$\underset{\theta_0, \theta_1}{\text{minimize}} J(\theta_0, \theta_1)$$


Cost Function

➤ Example

When $\theta_0 = 0$, and $\theta_1 = 0.5$

$$J(\theta_0, \theta_1) = \frac{1}{2n} \sum_{k=1}^n (0.5x^{(k)} - y^{(k)})^2 =$$




$$\frac{1}{2 \cdot 3} ((0.5 - 1)^2 + (1 - 2)^2 + (1.5 - 3)^2) = 0.58$$



When $\theta_0 = 0$, and $\theta_1 = 1$


$$J(\theta_0, \theta_1) = \frac{1}{2n} \sum_{k=1}^n (h_\theta(x^{(k)}) - y^{(k)})^2$$

$$\frac{1}{2n} \sum_{k=1}^n (x^{(k)} - y^{(k)})^2 = \frac{1}{2 \cdot 3} (0^2 + 0^2 + 0^2) = 0$$



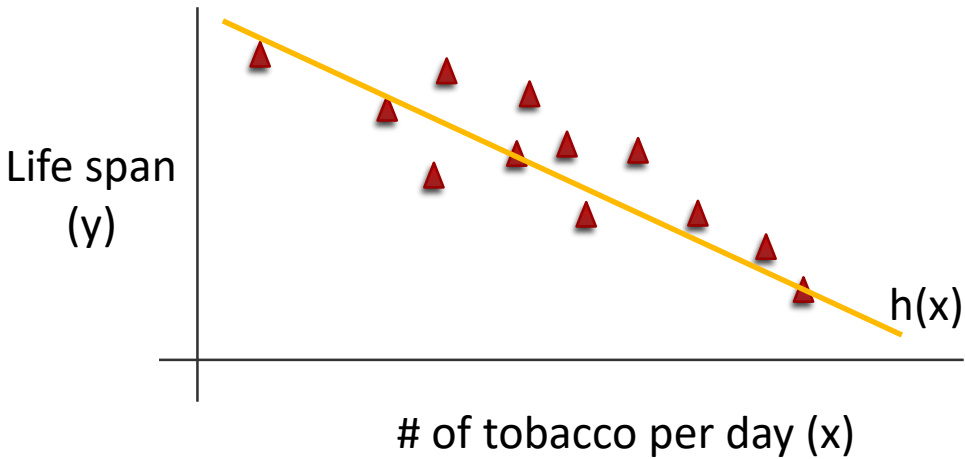
Cost Function

➤ $J(\theta_0, \theta_1)$ with two parameters




Cost Function


$$h_{\theta}(x) = \theta_0 + \theta_1 x$$



$$J(\theta_0, \theta_1)$$



Gradient Descent Algorithm




Gradient Descent Algorithm


➤ Repeat the function below until it converges

$$\theta_j := \theta_j - \alpha \frac{d}{d\theta_j} J(\theta_0, \theta_1) \text{ for } j=0 \text{ and } j=1$$

where $\alpha (>0)$ is a learning rate




$$\theta_1 := \theta_1 - \alpha \frac{d}{d\theta_1} J(\theta_1)$$




Gradient Descent Algorithm

$$\theta_1 := \theta_1 - \alpha \frac{d}{d\theta_1} J(\theta_1)$$

- Too small α makes gradient descent slow



- Too large α makes gradient descent could overshoot the minimum.




- Gradient descent can converge to a local minimum

Gradient Descent Algorithm

$$\theta_1 := \theta_1 - \alpha \frac{d}{d\theta_1} J(\theta_1)$$

- Since gradient descent takes smaller steps by itself, we don't have to adjust or decrease α over time.



Gradient Descent for Linear Regression

- Linear regression hypothesis

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

- Cost function

$$J(\theta_0, \theta_1) = \frac{1}{2n} \sum_{k=1}^n (h_{\theta}(x^{(k)}) - y^{(k)})^2$$

- Gradient descent algorithm

Repeat the function below until it converges

$$\theta_j := \theta_j - \alpha \frac{d}{d\theta_j} J(\theta_0, \theta_1) \text{ for } j=0 \text{ and } j=1$$

Gradient Descent for Linear Regression

$$\begin{aligned}\frac{d}{d\theta_j} J(\theta_0, \theta_1) &= \frac{d}{d\theta_j} \frac{1}{2n} \sum_{k=1}^n (h_{\theta}(x^{(k)}) - y^{(k)})^2 \\ &= \frac{d}{d\theta_j} \frac{1}{2n} \sum_{k=1}^n (\theta_0 + \theta_1 x^{(k)} - y^{(k)})^2\end{aligned}$$

$$j = 0 \Rightarrow \frac{d}{d\theta_0} J(\theta_0, \theta_1) = \frac{1}{n} \sum_{k=1}^n (h_{\theta}(x^{(k)}) - y^{(k)})$$

$$j = 1 \Rightarrow \frac{d}{d\theta_1} J(\theta_0, \theta_1) = \frac{1}{n} \sum_{k=1}^n (h_{\theta}(x^{(k)}) - y^{(k)}) x^{(k)}$$


Gradient Descent for Linear Regression

Repeat the function below until it converges


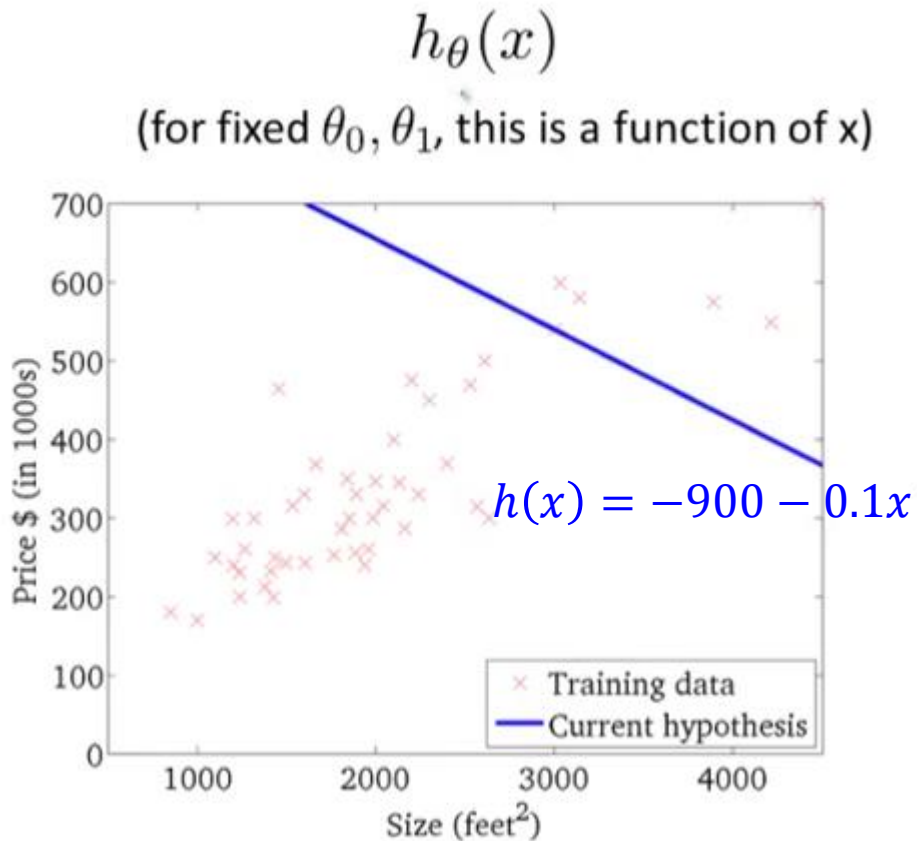
$$\theta_0 := \theta_0 - \alpha \frac{1}{n} \sum_{k=1}^n (h_{\theta}(x^{(k)}) - y^{(k)})$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{n} \sum_{k=1}^n (h_{\theta}(x^{(k)}) - y^{(k)}) x^{(k)}$$

Update θ_0 and θ_1 simultaneously



Gradient Descent for Linear Regression




Machine Learning by Andrew Ng.

Gradient Descent for Linear Regression


$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)



$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)




Machine Learning by Andrew Ng.

Gradient Descent for Linear Regression


$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)



$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)




Machine Learning by Andrew Ng.

Gradient Descent for Linear Regression


$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)



$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)




Machine Learning by Andrew Ng.

Gradient Descent for Linear Regression


$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)



$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)




Machine Learning by Andrew Ng.

Gradient Descent for Linear Regression


$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)



$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)




Machine Learning by Andrew Ng.

Gradient Descent for Linear Regression


$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)



$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)




Machine Learning by Andrew Ng.

Gradient Descent for Linear Regression


$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)



$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)




Machine Learning by Andrew Ng.

Gradient Descent for Linear Regression


$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)



$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)



Machine Learning by Andrew Ng.

Linear Regression with multiple variables

- When we have multiple variables(features) to interpret a phenomenon
 - for example, lifespan vs smoking/exercise/fine dust/diet and so on....

- Hypothesis for the multiple variables

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 + \theta_4 x_4 + \dots$$

$$\rightarrow h_{\theta}(x) = \theta^T X$$

$$, \text{ where } \theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \dots \\ \theta_n \end{bmatrix} \text{ and } X = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ \dots \\ x_n \end{bmatrix} \quad (x_0 = 1)$$

Linear Regression with multiple variables

- Hypothesis of linear regression

$$h_{\theta}(x) = \boldsymbol{\theta}^T x = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

- Cost function

$$J(\theta_0, \theta_1, \dots, \theta_n) = \frac{1}{2m} \sum_{k=1}^m (h_{\theta}(x^{(k)}) - y^{(k)})^2$$

- Gradient descent algorithm

Repeat the function below until it converges

$$\theta_j := \theta_j - \alpha \frac{d}{d\theta_j} J(\theta_0, \theta_1, \dots, \theta_n) \text{ for } j=0, \dots, n$$

Linear Regression with multiple variables

Repeat the function below until it converges

$$\theta_j := \theta_j - \alpha \frac{d}{d\theta_j} J(\theta_0, \theta_1, \dots, \theta_n) \text{ for } j=0, \dots, n$$

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{k=1}^m (h_{\theta}(x^{(k)}) - y^{(k)}) x_0^{(k)}$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{k=1}^m (h_{\theta}(x^{(k)}) - y^{(k)}) x_1^{(k)}$$

$$\theta_2 := \theta_2 - \alpha \frac{1}{m} \sum_{k=1}^m (h_{\theta}(x^{(k)}) - y^{(k)}) x_2^{(k)}$$

⋮

Feature Scaling

➤ Features should be in similar range

➤ Set the range into $-1 \leq x_i \leq 1$

➤ Feature normalization : zero mean and unit variance



$$\frac{x_i - \mu}{\sigma}$$

For example, $x_1 = \# \text{ of cigaret (0~40 pieces)}$

$x_2 = \text{seconds of exercise (0~7200 secs)}$


$$x_1 = \frac{\# \text{ of cigaret}}{40}$$

$$x_2 = \frac{\text{seconds of exercise}}{7200}$$



Convergence of Gradient Descent

- Checking the convergence of gradient descent




- Sometimes, the stopping criterion of gradient descent is set as

$$J(\theta) < 10^{-3}$$

- Using too large α , $J(\theta)$ may not decrease on every iteration


- Small α makes $J(\theta)$ decrease on every iteration

- However, too small α causes too slow convergence



Polynomial Regression

➤ Quadratic function




$$h_{\theta}(x) = \theta_0 + \theta_1 x_1$$

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2^2$$

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2^2 + \theta_3 x_3^3$$

* Square root function



$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 \sqrt{x_2}$$

Normal Equation

➤ Normal equation : analytic way to solve for θ

	Smoking (cigarette #)	Exercise (Seconds/day)	Fine dust ($\mu g/m^3$)	Diet (kcal)	Lifespan (years)
x_0	x_1	x_2	x_3	x_4	y
1	5	3600	20	2000	90
1	0	1000	60	2500	75
1	10	7200	150	3500	57
1	40	500	100	1000	45

$$\mathbf{X} = \begin{bmatrix} 1 & 5 & 3600 & 20 & 2000 \\ 1 & 0 & 1000 & 60 & 2500 \\ 1 & 10 & 7200 & 150 & 3500 \\ 1 & 40 & 500 & 100 & 1000 \end{bmatrix}$$

$$\mathbf{y} = \begin{bmatrix} 90 \\ 75 \\ 57 \\ 45 \end{bmatrix}$$

$$y = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 + \theta_4 x_4 \quad \Rightarrow \quad \theta = (X^T X)^{-1} X^T y$$

Gradient Descent vs Normal Equation

- When m is the sample numbers and n is the number of features
- Gradient Descent
 - Needs α
 - Takes many iterations
 - is OK with large n
- Normal Equation
 - Needs no α
 - Needs no iteration
 - Need to compute $(X^T X)^{-1}$
 - Slow with large n