

Lecture 8-1 Perceptron

Machine Learning

Spring Semester '2022


2nd Half Course Introduction

- Instructor: Hyukjoon Lee
hlee@kw.ac.kr
새빛관 707
Tel: 940-5127
Office Hours: T,W 2-3 pm
- Class Hour: T 4:30–5:45 pm (recorded video), Th 3:00–4:15 pm
- Class Room: 새빛관 202
- Course Homepage: mclab.kw.ac.kr
- Required Text:
 - Lecture slides
- Recommended References:
 - Ian Goodfellow, Yoshua Bengio, Aaron Courville, Deep Learning, Springer, 2016


Machine learning overview

 x $f(x)$ y


```
Hello,  
  
Do you want free printer  
cartridges? Why pay more  
when you can get them  
ABSOLUTELY FREE! Just
```




```
# free      : 2  
YOUR_NAME   : 0  
MISPELLED   : 2  
FROM_FRIEND : 0  
...
```



SPAM
or
+




```
PIXEL-7,12  : 1  
PIXEL-7,13  : 0  
...  
NUM_LOOPS   : 1  
...
```



"2"

Biological Model

- Human neurons




Linear classifiers

- Inputs are **feature values**
- Each feature has a **weight**
- Sum is the **activation**


$$\text{activation}_w(x) = \sum_i w_i \cdot f_i(x) = w \cdot f(x)$$

- If the activation is:
 - Positive, output +1
 - Negative, output -1



Weights

- Binary case: compare features to a weight vector
- Learning: figure out the weight vector from examples





Binary decision rule

- In the space of feature vectors
 - Examples are points
 - Any weight vector is a hyperplane
 - One side corresponds to $Y=+1$
 - Other corresponds to $Y=-1$


w

BIAS	:	-3
free	:	4
money	:	2
...		



Learning: binary perceptron

- Start with weights = 0
- For each training instance:
 - Classify with current weights
- If correct (i.e., $y=y^*$), no change!
- If wrong: adjust the weight vector




Learning: binary perceptron

- Start with weights = 0
- For each training instance:
 - Classify with current weights

$$y = \begin{cases} +1 & \text{if } w \cdot f(x) \geq 0 \\ -1 & \text{if } w \cdot f(x) < 0 \end{cases}$$

- If correct (i.e., $y=y^*$), no change!
- If wrong: adjust the weight vector by adding or subtracting the feature vector. Subtract if y^* is -1.

$$w = w + y^* \cdot f$$



Multiclass decision rule

- If we have multiple classes:
 - A weight vector for each class:


$$w_y$$

- Score (activation) of a class y :

$$w_y \cdot f(x)$$

- Prediction highest score wins

$$y = \arg \max_y w_y \cdot f(x)$$



Binary = multiclass where the negative class has weight zero

Learning: multiclass perceptron


- Start with all weights = 0
- Pick up training examples one by one
- Predict with current weights

$$y = \arg \max_y w_y \cdot f(x)$$

- If correct, no change!
- If wrong: lower score of wrong answer, raise score of right answer

$$w_y = w_y - f(x)$$

$$w_{y^*} = w_{y^*} + f(x)$$



Example: multiclass perceptron (1)

“win the vote”

“win the election”

“win the game”

w_{SPORTS}

BIAS	:	1
win	:	0
game	:	0
vote	:	0
the	:	0
...		

$w_{POLITICS}$

BIAS	:	0
win	:	0
game	:	0
vote	:	0
the	:	0
...		

w_{TECH}

BIAS	:	0
win	:	0
game	:	0
vote	:	0
the	:	0
...		

Example: multiclass perceptron (2)

$$w_s: [1 \ 0 \ 0 \ 0 \ 0]^T \quad w_p: [0 \ 0 \ 0 \ 0 \ 0]^T \quad w_t: [0 \ 0 \ 0 \ 0 \ 0]^T$$

$$x_1: \text{"win the vote"} \rightarrow [1 \ 1 \ 0 \ 1 \ 1]^T \quad y = \text{'politics'} \ (p)$$

$$x_2: \text{"win the election"} \rightarrow [1 \ 1 \ 0 \ 0 \ 1]^T \quad y = \text{'politics'} \ (p)$$

$$x_3: \text{"win the game"} \rightarrow [1 \ 1 \ 1 \ 0 \ 1]^T \quad y = \text{'sports'} \ (s)$$

$$w_s \cdot x_1 = 1, \quad w_p \cdot x_1 = 0, \quad w_t \cdot x_1 = 0, \quad y = s, \quad y^* = p \rightarrow \text{incorrect!}$$

$$\rightarrow w_p = [0 \ 0 \ 0 \ 0 \ 0]^T + [1 \ 1 \ 0 \ 1 \ 1]^T = [1 \ 1 \ 0 \ 1 \ 1]^T$$

$$\rightarrow w_s = [1 \ 0 \ 0 \ 0 \ 0]^T - [1 \ 1 \ 0 \ 1 \ 1]^T = [0 \ -1 \ 0 \ -1 \ -1]^T$$


$$w_s \cdot x_2 = 1$$

Properties of perceptrons


- Separability: true if some parameters get the training set perfectly correct
- Convergence: if the training is separable, perceptron will eventually converge (binary case)
- Mistake Bound: the maximum number of mistakes (binary case) related to the *margin* or degree of separability

$$\text{mistakes} < \frac{k}{\delta^2}$$

Separable




Non-Separable




Examples: perceptron

- Non-Separable Case




Problems with the perceptron

- Noise: if the data isn't separable, weights might thrash
 - Averaging weight vectors over time can help (averaged perceptron)



- Mediocre generalization: finds a “barely” separating solution



- Overtraining: test / held-out accuracy usually rises, then falls
 - Overtraining is a kind of overfitting

