

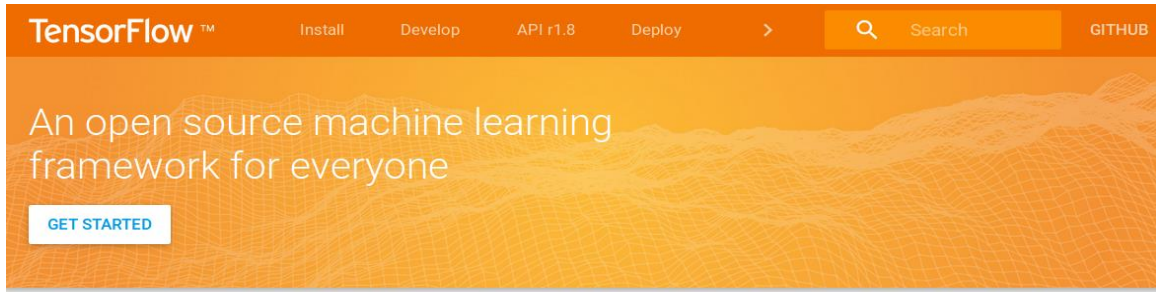
Lecture 12 – TensorFlow Basics

Machine Learning

TensorFlow Basics

1. **Why Tensorflow?**
2. **Tf Website**
3. **Data Flow Graph**
4. **Installation And Simple Validation**
5. **Say Hello For TF With Jupyter Notebook**
6. **TF Mechanics And Computational Graph**
7. **Tensor, Placeholder And Variable**
8. **MNIST Dataset and its TF Example**

TF Website: <https://www.tensorflow.org>



TensorFlow 1.8 has arrived!

We're excited to announce the release of TensorFlow 1.8! Check out the announcement to upgrade your code with ease.

[LEARN MORE](#)



TensorFlow Dev Summit 2018

Thousands of people from the TensorFlow community participated in the second TensorFlow Dev Summit. Watch the keynote and talks now.

[WATCH NOW](#)



Announcing TensorFlow.js!

Learn more about our new library for machine learning in the browser using JavaScript.

[LEARN MORE](#)

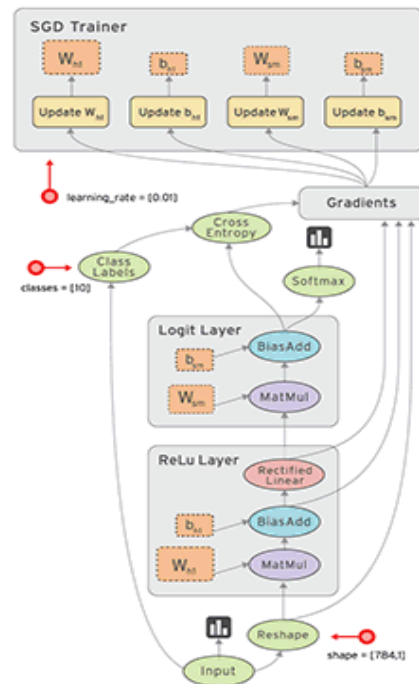
TensorFlow

- TensorFlow™ is an open source software library for numerical computation using data flow graphs.
- But what does it actually do?
 - TensorFlow provides primitives for defining functions on tensors and automatically computing their derivatives.
- With Python!



What is a Data Flow Graph?

- **Dataflow** is a common programming model for parallel computing
 - the *nodes* represent units of computation,
 - the *edges* represent the data consumed or produced by a computation.
- **What' s a Tensor?**
 - Matrix, vector, scalar
 - Formally, tensors are multilinear maps from vector spaces to the real numbers
 - a tensor can be represented as a multidimensional array of numbers.



Everything is Tensor

Tensors

```
In [3]: 3 # a rank 0 tensor; this is a scalar with shape []  
[1., 2., 3.] # a rank 1 tensor; this is a vector with shape [3]  
[[1., 2., 3.], [4., 5., 6.]] # a rank 2 tensor; a matrix with shape [2, 3]  
[[[1., 2., 3.]], [[7., 8., 9.]]] # a rank 3 tensor with shape [2, 1, 3]
```

- The central unit of data in TensorFlow is the **tensor**.
- A tensor consists of a set of primitive values shaped into an array of any number of dimensions.
- A tensor's **rank** is its number of dimensions,
- while its **shape** is a tuple of integers specifying the array's length along each dimension
- TensorFlow uses numpy arrays to represent tensor **values**.

Tensor Ranks, Shapes, and Types

```
t = [[1, 2, 3], [4, 5, 6], [7, 8, 9]]
```

Rank	Math entity	Python example
0	Scalar (magnitude only)	s = 483
1	Vector (magnitude and direction)	v = [1.1, 2.2, 3.3]
2	Matrix (table of numbers)	m = [[1, 2, 3], [4, 5, 6], [7, 8, 9]]
3	3-Tensor (cube of numbers)	t = [[[2], [4], [6]], [[8], [10], [12]], [[14], [16], [18]]]
n	n-Tensor (you get the idea)

Rank	Shape	Dimension number	Example
0	[]	0-D	A 0-D tensor. A scalar.
1	[D0]	1-D	A 1-D tensor with shape [5].
2	[D0, D1]	2-D	A 2-D tensor with shape [3, 4].
3	[D0, D1, D2]	3-D	A 3-D tensor with shape [1, 4, 3].
n	[D0, D1, ... Dn-1]	n-D	A tensor with shape [D0, D1, ... Dn-1].

Jupyter notebook

- Jupyter is an web-based interactive development environment.
- It support multiple programming languages like Python.
 - When you install the Anaconda,
 - you get your own jupyter notebook at the same time!
- For more details
 - <https://jupyter.readthedocs.io/en/latest/index.html>
 - Video Tutorial: <https://www.youtube.com/user/roshanRush>

How to run?

`$ jupyter notebook`



TensorFlow Hello World!

Hello TensorFlow!

```
In [2]: # Create a constant op
# This op is added as a node to the default graph
hello = tf.constant("Hello, TensorFlow!")

# start a TF session
sess = tf.Session()

# run the op and get result
print(sess.run(hello))

b'Hello, TensorFlow!'
```

'b' indicates Bytes literals.


TensorFlow Mechanics

- 1 Build **graph** using TensorFlow **operations**

tf.Graph

- 2 **feed data** and **run graph**
(**Session** -> **operations**)
sess.run (ops)

- 3 **update variables** in the **graph**
(and **return values**)




Computational Graph


```
In [4]: node1 = tf.constant(3.0, tf.float32)
        node2 = tf.constant(4.0) # also tf.float32 implicitly
        node3 = tf.add(node1, node2)
```

```
In [5]: print("node1: ", node1)
        print("node2: ", node2)
        print("node3: ", node3)
```

```
node1: Tensor("Const_1:0", shape=(), dtype=float32)
node2: Tensor("Const_2:0", shape=(), dtype=float32)
node3: Tensor("Add:0", shape=(), dtype=float32)
```



- A **computational graph** is a series of TensorFlow **operations** arranged into a **graph**.
- The **graph** is composed of two types of objects.
 - **Operations** (or "**ops**"): The nodes of the graph. Operations describe calculations that *consume* and *produce* **tensors**.
 - **Tensors**: The edges in the graph. These represent the values that will flow through the graph. Most TensorFlow functions return `tf.Tensors`.




Computational Graph

- (1) Build graph (tensors) using TensorFlow operations

```
In [4]: node1 = tf.constant(3.0, tf.float32)
        node2 = tf.constant(4.0) # also tf.float32 implicitly
        node3 = tf.add(node1, node2)
```

- (2) feed data and run graph (operation)
sess.run (op)
- (3) update variables in the graph (and return values)



Computational Graph [Session]

- (2) feed data and run graph (operation)
sess.run (op)
- (3) update variables in the graph (and return values)

```
In [6]: sess = tf.Session()
print("sess.run(node1, node2): ", sess.run([node1, node2]))
print("sess.run(node3): ", sess.run(node3))

sess.run(node1, node2): [3.0, 4.0]
sess.run(node3): 7.0
```

create a simple session

```
sess = tf.Session()
```

...

```
sess.close()
```

Or Create a default in-process session for the current default graph.

```
with tf.Session() as sess:
```

...

Computational Graph

▪ Session

- To evaluate tensors, instantiate a `tf.Session` object, informally known as a `sess(ion)`.
- A session encapsulates the state of the TensorFlow runtime, and `runs` TensorFlow `operations`.
 - ✓ If a `tf.Graph` is like a .py file, a `tf.Session` is like the python executable.

- You can pass *multiple tensors* to `tf.Session.run`.

The run method transparently handles any combination of tuples or dictionaries

- During a call to `tf.Session.run` any `tf.Tensor` only has a `single` value

```
In [5]: vec = tf.random_uniform(shape=(3,))
out1 = vec + 1
out2 = vec + 2
print(sess.run(vec))
print(sess.run(vec))
print(sess.run(out1))
print(sess.run(out2))
print(sess.run((out1, out2)))

[0.18464649 0.14886129 0.5047021 ]
[0.3270713  0.4014722  0.22448409]
[1.3578813  1.7317374  1.8207817]
[2.3646321  2.3802166  2.8418117]
(array([1.3722222, 1.2407451, 1.8812546], dtype=float32),
 array([2.3722222, 2.240745 , 2.8812547], dtype=float32))
```

The result shows a **different** random value **on each call** to run,

but a **consistent** value during a **single run**

(out1, out2) receive the same random input

Placeholders and Feeding

```
In [7]: a = tf.placeholder(tf.float32)
        b = tf.placeholder(tf.float32)
        adder_node = a + b  # + provides a shortcut for tf.add(a, b)

        print(sess.run(adder_node, feed_dict={a: 3, b: 4.5}))
        print(sess.run(adder_node, feed_dict={a: [1,3], b: [2, 4]}))

7.5
[ 3.  7.]
```

- A graph can be parameterized to accept external inputs, known as **placeholders**
- A **placeholder** is a promise to provide a value later, like a function argument.
- We can evaluate this graph with multiple inputs by using the **feed_dict** argument of the run method to feed concrete values to the **placeholders**

Placeholders and Feeding (without `feed_dict` and add closed session)

```
In [10]: import tensorflow as tf
a = tf.placeholder(tf.float32)
b = tf.placeholder(tf.float32)
adder_node = a + b # + provides a shortcut for tf.add(a, b)
sess = tf.Session()
print(sess.run(adder_node, {a: 3, b: 4.5}))
sess.close()
print(sess.run(adder_node, {a: 3, b: 4.5}))
```

7.5

```
-----
RuntimeError                                Traceback (most recent call l
ast)
<ipython-input-10-2ecc9be8ffc4> in <module>()
      6 print(sess.run(adder_node, {a: 3, b: 4.5}))
      7 sess.close()
----> 8 print(sess.run(adder_node, {a: 3, b: 4.5}))
~/anaconda3/envs/tensorflow/lib/python3.6/site-packages/tensorflow/pyth
on/client/session.py in _run(self, handle, fetches, feed_dict, options,
run_metadata)
    1056     # Check session.
    1057     if self._closed:
-> 1058         raise RuntimeError('Attempted to use a closed Session.')
    1059     if self.graph.version == 0:
    1060         raise RuntimeError('The Session graph is empty. Add oper
ations to the '
RuntimeError: Attempted to use a closed Session.
```


Variables

- A TensorFlow **variable** is the best way to represent shared, persistent state manipulated by your program.
- *The best way to create a variable is to call the **tf.get_variable()** function.*

```
In [1]: import tensorflow as tf
```



```
In [2]: # simply provide the name and shape
my_variable = tf.get_variable("my_variable", [1, 2, 3])
```

This variable will, by **default**, have the **dtype** `tf.float32` and its initial value will be **randomized** via `tf.glorot_uniform_initializer`.

```
In [3]: # specify the dtype and initializer
my_int_variable = tf.get_variable("my_int_variable", [1, 2, 3],
                                   dtype=tf.int32, initializer=tf.zeros_initializer)

# initialize a tf.Variable to have the value of a tf.Tensor.
other_variable = tf.get_variable("other_variable",
                                   dtype=tf.int32, initializer=tf.constant([23, 42]))
```



MNIST Dataset



[train-images-idx3-ubyte.gz](#): training set images (9912422 bytes)
[train-labels-idx1-ubyte.gz](#): training set labels (28881 bytes)
[t10k-images-idx3-ubyte.gz](#): test set images (1648877 bytes)
[t10k-labels-idx1-ubyte.gz](#): test set labels (4542 bytes)


MNIST Dataset

- Check out https://www.tensorflow.org/get_started/mnist/beginners for more information about the mnist dataset
- MNIST data image of shape $28 * 28 = 784$
- $X = \text{tf.placeholder}(\text{tf.float32}, [\text{None}, 784])$
- $\# 0 - 9 \text{ digits recognition} = 10 \text{ classes}$
- $Y = \text{tf.placeholder}(\text{tf.float32}, [\text{None}, \text{nb_classes}])$



MNIST Dataset & EXAMPLE

```
import tensorflow
```



```
In [1]: import tensorflow as tf
import numpy as np
```

```
In [2]: from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)
```

Dataset	Purpose
<code>input_data.train</code>	55000 images and labels, for primary training.
<code>input_data.validation</code>	5000 images and labels, for iterative validation of training accuracy.
<code>input_data.test</code>	10000 images and labels, for final testing of trained accuracy.

MNIST Dataset

get the data

```
In [2]: # if there is no "MNIST_data/" then download
from tensorflow.examples.tutorials.mnist import input_data
# extract and read the mnist data
mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)
```

Successfully downloaded train-images-idx3-ubyte.gz 9912422 bytes.
Extracting MNIST_data/train-images-idx3-ubyte.gz
Successfully downloaded train-labels-idx1-ubyte.gz 28881 bytes.
Extracting MNIST_data/train-labels-idx1-ubyte.gz
Successfully downloaded t10k-images-idx3-ubyte.gz 1648877 bytes.
Extracting MNIST_data/t10k-images-idx3-ubyte.gz
Successfully downloaded t10k-labels-idx1-ubyte.gz 4542 bytes.
Extracting MNIST_data/t10k-labels-idx1-ubyte.gz

one-hot

```
In [48]: one_hot = tf.one_hot([[0],[1],[2],[0]],depth=3)

In [49]: sess.run(tf.one_hot([[0],[1],[2],[0]],depth=3))
Out[49]: array([[[ 1.,  0.,  0.],
                  [ 0.,  1.,  0.],
                  [ 0.,  0.,  1.],
                  [ 1.,  0.,  0.]], dtype=float32)

In [50]: sess.run(tf.one_hot([[0],[1],[2],[0]],depth=4))
Out[50]: array([[[ 1.,  0.,  0.,  0.],
                  [ 0.,  1.,  0.,  0.],
                  [ 0.,  0.,  1.,  0.],
                  [ 1.,  0.,  0.,  0.]], dtype=float32)

In [51]: sess.run(tf.reshape(one_hot,shape=[-1,3]))
Out[51]: array([[ 1.,  0.,  0.],
                [ 0.,  1.,  0.],
                [ 0.,  0.,  1.],
                [ 1.,  0.,  0.]], dtype=float32)
```

If you have download the MNIST dataset

Extracting MNIST_data/train-images-idx3-ubyte.gz
Extracting MNIST_data/train-labels-idx1-ubyte.gz
Extracting MNIST_data/t10k-images-idx3-ubyte.gz
Extracting MNIST_data/t10k-labels-idx1-ubyte.gz

MNIST EXAMPLE

set variables

```
In [4]: nb_class = 10

# MNIST data image of shape 28 * 28 = 784
X = tf.placeholder(tf.float32,[None,784])
# 0-9 digits recognition = 10 classes
Y = tf.placeholder(tf.float32,[None, nb_class])

W = tf.get_variable("w", [784,nb_class],
                    initializer=tf.truncated_normal_initializer(stddev=0.1))
b = tf.get_variable("biases", [nb_class],
                    initializer=tf.constant_initializer(0.1))
```

The *None* means, the value will be got by $X / 784$

MNIST EXAMPLE

softmax for the linear regression

```
In [4]: # Hypothesis (using softmax)
hypothesis = tf.nn.softmax(tf.matmul(X,W)+b)

cost = tf.reduce_mean(-tf.reduce_sum(Y * tf.log(hypothesis),axis = 1))
optimizer = tf.train.GradientDescentOptimizer(
    learning_rate=0.1).minimize(cost)

# test model
is_correct = tf.equal(tf.argmax(hypothesis,1),tf.argmax(Y,1))
# Caculate accuracy
accuracy = tf.reduce_mean(tf.cast(is_correct,tf.float32))
```

nn is a module and *softmax* is an operation(function)

MNIST EXAMPLE

training epoch/batch

```
In [5]: # parameters
training_epochs = 15
batch_size = 100

with tf.Session() as sess:
    # initialize Tensorflow Variables
    sess.run(tf.global_variables_initializer())
    # training cycles
    for epoch in range(training_epochs):
        avg_cost = 0
        total_batch = int(mnist.train.num_examples/batch_size)

        for i in range(total_batch):
            batch_xs,batch_ys = mnist.train.next_batch(batch_size)
            c,_ = sess.run([cost,optimizer],feed_dict={X:batch_xs,Y:batch_ys})
            avg_cost += c / total_batch

        print("Epoch:", '%04d'%(epoch+1)
              , 'cost=', '{:.9f}'.format(avg_cost))

    print("Learning finished")

    # Test the model using test sets
    print("old Accuracy: ", accuracy.eval(session=sess, feed_dict={
        X: mnist.test.images, Y: mnist.test.labels}))
    print('new Accuracy:', sess.run(accuracy, feed_dict={
        X: mnist.test.images, Y: mnist.test.labels}))
```


MNIST EXAMPLE

```
# Sample image show and prediction
import matplotlib.pyplot as plt
import random

# Get one and predict
r = random.randint(0,mnist.test.num_examples -1 )
print("Label:",sess.run(tf.argmax(mnist.test.labels[r:r+1],1)))
print("Prediction:",sess.run(tf.argmax(hypothesis,1),
                                feed_dict={X:mnist.test.images[r:r+1]}))

plt.imshow(mnist.test.images[r:r+1].reshape(28,28),
            cmap='Greys',interpolation='nearest')
plt.show()
```

MNIST EXAMPLE

```
Epoch: 0001 cost= 2.804654121
Epoch: 0002 cost= 1.179020686
Epoch: 0003 cost= 0.921789622
Epoch: 0004 cost= 0.797202877
Epoch: 0005 cost= 0.720808540
Epoch: 0006 cost= 0.666504125
Epoch: 0007 cost= 0.625679149
Epoch: 0008 cost= 0.593117082
Epoch: 0009 cost= 0.566464136
Epoch: 0010 cost= 0.544340848
Epoch: 0011 cost= 0.525419921
Epoch: 0012 cost= 0.508744514
Epoch: 0013 cost= 0.494571513
Epoch: 0014 cost= 0.481707098
Epoch: 0015 cost= 0.470132457
Learning finished
Accuracy: 0.8881
Label: [7]
Prediction: [7]
```

